Hongseok Namkoong

Columbia Business School
Kravis 933, 625 W 130th Street
New York, NY 10027

Email: namkoong@gsb.columbia.edu
Homepage: http://hsnamkoong.github.io
Google Scholar: https://bit.ly/3hmDmjY

Employment

Assistant Professor, Decision, Risk, and Operations Division, Columbia Business School, 2020—Present

LinkedIn Scholar, Trust & Responsible AI, LinkedIn, 2022–Present

Research Scientist, Core Data Science, Facebook, 2019–2020

Research Assistant, Peter W. Glynn and John C. Duchi, Stanford University, 2014–2019

Education

Ph.D. Management Science and Engineering, Stanford University, 2019
Advisors: John C. Duchi and Peter W. Glynn

M.S. Statistics, Stanford University, 2017

B.S. Summa Cum Laude. Industrial Engineering and Mathematics, KAIST, 2013

Teaching

B8103: Business Analytics II (MBA, MS), Columbia University, 2021-2023

B9145: Reliable Statistical Learning (PhD), Columbia University, 2020, 2023

Publications[1]

Preprints

P1 A. Boyarsky, H. Namkoong, and J. Pouget-Abadie. Modeling interference via experiment rollout. *arXiv:2305.10728 [stat.ME]*, 2023. Conference version appeared in ACM conference on Economics and Computation.

P2 T. Cai, H. Namkoong, and S. Yadlowsky. Diagnosing model performance under distribution shift. *arXiv:2303.02011 [stat.ML]*, 2023. Conference version appeared Symposium on Foundations of Responsible Computing 2023

P3 H. Namkoong, Y. Ma, and P. W. Glynn. Minimax optimal estimation of stability under distribution shift. *arXiv:2212.06338 [stat.ML]*, 2022. Major revision in Operations Research.

P4 M. Li, H. Namkoong, and S. Xia. Evaluating model performance under worst-case subpopulations. *Working paper*, 2022. Short version appeared in NeurIPS 2021.

P5 H. Namkoong, S. Daulton, and E. Bakshy. Distilled thompson sampling: Practical and efficient thompson sampling via imitation learning. *arXiv:2011.14266 [cs.LG]*, 2021. Selected for an oral presentation at the Neurips 2020 OfflineRL Workshop.

---

[1]Customary authorship ordering is by alphabetical order. Name* denotes equal contribution.

P6 S. Jeong and H. Namkoong. Robust causal inference under covariate shift via worst-case subpopulation treatment effect. *arXiv:2007.02411 [stat.ML]*, 2022. Short version appeared in Conference on Learning Theory 2020.

Journal publications

J1 S. Yadlowsky, H. Namkoong, S. Basu, J. Duchi, and L. Tian. Bounds on the conditional and average treatment effect with unobserved confounding factors. *Annals of Statistics*, 50(5):2587–2615, 2022.

J2 J. C. Duchi, T. Hashimoto, and H. Namkoong. Distributionally robust losses against mixture covariate shifts. *Operations Research*, 2022.

J3 J. C. Duchi and H. Namkoong. Learning models with uniform performance via distributionally robust optimization. *Annals of Statistics*, 49(3):1378–1406, 2021.

J4 J. C. Duchi, P. W. Glynn, and H. Namkoong. Statistics of robust optimization: A generalized empirical likelihood approach. *Mathematics of Operations Research*, 46(3):946–969, 2021.

J5 J. C. Duchi and H. Namkoong. Variance-based regularization with convex objectives. *Journal of Machine Learning Research*, 2019.

Refereed conference proceedings[2]

C1 A. Boyarsky, H. Namkoong, and J. Pouget-Abadie. Modeling interference via experiment rollout. In *Proceedings of the 24th ACM conference on Economics and Computation*, 2023.

C2 T. Cai, H. Namkoong, and S. Yadlowsky. Diagnosing model performance under distribution shift. In *Symposium on Foundations of Responsible Computing*, 2023.

C3 M. Wortsman, G. Ilharco, S. Y. Gadre, R. Roelofs, R. Gontijo-Lopes, A. S. Morcos, H. Namkoong, A. Farhadi, Y. Carmon, S. Kornblith, and L. Schmidt. Model soups: averaging weights of multiple fine-tuned models improves accuracy without increasing inference time. In *Proceedings of the 39th International Conference on Machine Learning*, 2022

C4 M. Wortsman, G. Ilharco, J. W. Kim, M. Li, S. Kornblith, R. Roelofs, R. Gontijo-Lopes, H. Hajishirzi, A. Farhadi, H. Namkoong, and L. Schmidt. Robust fine-tuning of zero-shot models. In *Proceedings of the 32nd IEEE Conference on Computer Vision and Pattern Recognition*, 2022. Selected for a full oral presentation

C5 M. Li, H. Namkoong, and S. Xia. Evaluating model performance under worst-case subpopulations. In *Advances in Neural Information Processing Systems 34*, 2021.

C6 H. Namkoong*, R. Keramati*, S. Yadlowsky*, and E. Brunskill. Off-policy policy evaluation for sequential decisions under unobserved confounding. In *Advances in Neural Information Processing Systems 33*, 2020.

C7 S. Jeong and H. Namkoong. Robust causal inference under covariate shift via worst-case subpopulation treatment effect. In *Conference on Learning Theory*, 2020.

C8 M. O'Kelly*, A. Sinha*, H. Namkoong*, J. Duchi, and R. Tedrake. Scalable end-to-end autonomous vehicle testing via rare-event simulation. In *Advances in Neural Information Processing Systems 31*, 2018.

C9 R. Volpi*, H. Namkoong*, J. Duchi, V. Murino, and S. Savarese. Generalizing to unseen domains via adversarial data augmentation. In *Advances in Neural Information Processing Systems 31*, 2018.

---

[2]Papers displayed in gray are superseded by long versions.

C10 T. Hashimoto, M. Srivastava, H. Namkoong, and P. Liang. Fairness without demographics in repeated loss minimization. In *International Conference on Machine Learning*, 2018. Best Paper Runner-up Award.

C11 A. Sinha*, H. Namkoong*, and J. Duchi. Certifiable distributional robustness with principled adversarial training. In *International Conference on Learning Representations*, 2018. Selected for a full oral presentation; 2% of submissions.

C12 H. Namkoong and J. C. Duchi. Variance regularization with convex objectives. In *Advances in Neural Information Processing Systems 30*, 2017. Best Paper Award.

C13 H. Namkoong, A. Sinha, S. Yadlowsky, and J. C. Duchi. Adaptive sampling probabilities for non-smooth optimization. In *International Conference on Machine Learning*, pages 2574–2583, 2017.

C14 H. Namkoong and J. C. Duchi. Stochastic gradient methods for distributionally robust optimization with f-divergences. In *Advances in Neural Information Processing Systems 29*, 2016.

Technical reports & software

T1 G. Ilharco, M. Wortsman, N. Carlini, R. Taori, A. Dave, V. Shankar, H. Namkoong, J. Miller, H. Hajishirzi, A. Farhadi, and L. Schmidt. Openclip: Open-sourced implementation of clip, Jul 2021. URL https://doi.org/10.5281/zenodo.5143773

T2 A. Sinha*, H. Namkoong*, R. Volpi, and J. Duchi. Certifying some distributional robustness with principled adversarial training. *Technical Report*, 2020.

T3 H. Namkoong, J. C. Duchi, and P. W. Glynn. Proofs for empirical likelihood with general f-divergences. *Technical Report*, 2018.

Honors & Awards

Amazon Research Award, 2022

Best Paper Finalist for "Robust fine-tuning of zero-shot models", *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022

Best Student Paper Award for "Statistics of Robust Optimization: A Generalized Empirical Likelihood Approach", *INFORMS Applied Probability Society*, 2018

Best Paper Runner Up Award for "Fairness Without Demographics in Repeated Loss Minimization", *International Conference on Machine Learning (ICML)*, 2018

Best Paper Award for "Variance Based Regularization with Convex Objectives", *Neural Information Processing Systems (NeurIPS)*, 2017

Samsung Fellowship, 2013–2018

Department Fellowship, Management Science and Engineering, Stanford, 2013–2018

KAIST President's Award (graduated top of class in the School of Engineering), 2013

Undergraduate Research Award, First Place, Department of Industrial and Systems Engineering, 2012

Professional Service

Senior Program Committee

| 2021-23 | Area chair, NeurIPS |
| 2023 | Area chair, ICML |

Reviewing

Journals *Operations Research, Management Science, Journal of the American Statistical Association, Journal of the Royal Statistical Society: Series B, Mathematical Programming, SIAM Journal on Mathematics of Data Science, Journal of Machine Learning Research, Transactions on Pattern Analysis and Machine Intelligence, Automatica.*

Conferences *Neural Information Processing Systems, International Conference on Machine Learning, Conference on Learning Theory, Conference on Algorithmic Learning Theory*

Workshop Organization

2023	Co-Organizer, three INFORMS invited sessions on experimentation
2023	Organizer, INFORMS APS invited session "Frontiers in Sequential Learning"
2022	Co-organizer, NeurIPS workshop on "Distribution shifts: connecting methods and applications (DistShift)"
2022	Organizer, INFORMS invited session "Causality & Robustness"
2022	Co-organizer, INFORMS invited session "Experimentation Design"
2022	INFORMS HAS Pierskalla Award committee member
2022	Co-organizer, ICDM workshop on "Algorithms Towards Ethical and Privacy Challenges in Social Media Recommendation Systems (AESM)"
2021	Organizer, mentoring program for PhD students at the Neurips workshop on "Distribution shifts: connecting methods and applications (DistShift)"
2021	Co-organizer, NeurIPS workshop on "Distribution shifts: connecting methods and applications (DistShift)"
2021	Co-organizer, JSM invited session on "distributional robustness, validity, causality, and generalizability"
2019	Co-organizer, INFORMS invited session on "AI and machine learning"

Research Supervision

Ph.D. Advising

Ari Boyarsky (DRO)

Tiffany Cai (Statistics)

Ethan Che (DRO, co-advised with Jing Dong)

Yuanzhe Ma (IEOR, co-advised with Garud Iyengar and Jay Sethuraman)

AYeong Lee (DRO)

Daksh Mittal (DRO)

Naimeng Ye (DRO)

Yibo Zeng (IEOR, co-advised with Henry Lam)

Ph.D. Thesis Committee

Nicholas Galbraith (Statistics), Yuanling (Judy) Gan (DRO), Yian Huang (Statistics), Shangzhou (Shawn) Xia (DRO), Yunbei Xu (DRO), Yibo Zeng (IEOR), Junzhe Zhang (CS)

Undergraduate Advising

Jimmy Wang (CS), Leon Li (CS), Jacklyn Tsai (CS), Rohan Subramani (CS), Martha Wangechi Njuguna (CS)

Invited Talks

2023	Joint Statistical Meetings, Toronto
2023	New England Statistics Symposium
2023	INFORMS APS Conference, Nancy, France
2023	Panelist, Showcase on Large Language Models, Columbia Center of AI Technology
2023	Department of Management Science and Engineering, Stanford University
2023	AI Matrix Seminar, UT San Antonio
2023	Moderator, Bernstein Center for Leadership and Ethics, Columbia Business School
2023	Operations Research Center, MIT
2022	Conference on Digital Experimentation, Cambridge
2022	INFORMS Annual Meeting, Indianapolis
2022	IFDS Workshop on Distributional Robustness in Data Science, University of Washington
2022	International Conference on Continuous Optimization (Bethlehem, PA)
2022	Department of Industrial & Systems Engineering and Computer Science, KAIST
2022	College of Business and AI, KAIST
2022	International Conference on Econometrics and Statistics, Kyoto
2022	Data Science Lab, MIT
2022	Department of Economics, Columbia University
2022	Workshop on Foundations of Stable, Generalizable and Transferable Statistical Learning, Mathematical Sciences Research Institute
2021	Department of Statistics, Columbia University
2021	LinkedIn
2021	Workshop on "distributional robustness, validity, causality, and generalizability", Joint Statistical Meetings
2021	Empirical Inference Department, Max Planck Institute for Intelligent Systems
2021	Department of Mathematics, KAIST
2021	School of Data Science, Seoul National University
2021	Data Science Institute, Columbia University
2021	Decision Science Group, McCombs School of Business, UT Austin
2020	Samsung Advanced Institute of Technology, Seoul
2020	Google Brain, Cambridge
2020	Cancelled due to COVID-19: Conference on Information Sciences and Systems, American Causal Inference Conference, SIAM Conference on Mathematics of Data Science
2019	Uber Marketplace and Uber Eats, San Francisco
2019	OIT Division, Graduate School of Business, Stanford University
2019	INFORMS Annual Meeting, Seattle
2019	Stitchfix, San Francisco
2019	Department of Computer Science, University of Wisconsin-Madison
2019	Department of Industrial and Systems Engineering, University of Wisconsin-Madison
2019	School of Operations Research and Industrial Engineering, Cornell Tech
2019	Machine Learning and Statistics Group, Microsoft Research New England
2019	Operations and Statistics Group, MIT Sloan School of Management
2019	Department of Operations Research and Industrial Engineering, UT Austin
2019	Machine Learning Department, Carnegie Mellon University
2019	Heinz College, Carnegie Mellon University

2019	Department of Industrial Engineering and Operations Research, Columbia University
2019	Decisions, Risk and Operations Division, Columbia Business School
2019	Department of Electrical and Computer Engineering, Purdue University
2019	Operations Management Division, Booth School of Business, University of Chicago
2019	Data Sciences and Operations, Marshall School of Business, University of Southern California
2018	Department of Industrial and Operations Engineering, University of Michigan
2018	Three invited talks, INFORMS Annual Meeting (Phoenix, AZ)
2018	Oral Presentation, International Conference on Learning Representations (Vancouver, Canada)
2017	Oral Presentation, Neural Information Processing Systems (Long Beach, CA)
2016	Department of Industrial and Systems Engineering, KAIST
2016	Young Researchers Workshop, School of ORIE, Cornell University

Academic Committees

Member, Ph.D. Admissions Committee, Columbia University, 2020–current

Member, Faculty Hiring Committee, Columbia University, 2021-2022

Member, Ph.D. Program Review Committee, Columbia University, 2021

Outside Activities   Columbia Business School requires faculty members to disclose any activities that might present a real or apparent conflict of interest. I consult at LinkedIn, serving as a LinkedIn Scholar.